


# Review: Pastry routing tables

|       |                                     |                         |                                     |                                     |                                     |                                     |                                     |                                     |                                     |                                     |  |  |  |  |  |  |
|-------|-------------------------------------|-------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|--|--|--|--|--|--|
| Row 0 | 0<br><i>x</i>                       | 1<br><i>x</i>           | 2<br><i>x</i>                       | 3<br><i>x</i>                       | 4<br><i>x</i>                       | 5<br><i>x</i>                       |                                     | 7<br><i>x</i>                       | 8<br><i>x</i>                       | 9<br><i>x</i>                       | <i>a</i><br><i>x</i>                       | <i>b</i><br><i>x</i>                       | <i>c</i><br><i>x</i>                       | <i>d</i><br><i>x</i>                       | <i>e</i><br><i>x</i>                       | <i>f</i><br><i>x</i>                       |
| Row 1 | 6<br>0<br><i>x</i>                  | 6<br>1<br><i>x</i>      | 6<br>2<br><i>x</i>                  | 6<br>3<br><i>x</i>                  | 6<br>4<br><i>x</i>                  |                                     | 6<br><i>x</i>                       | 6<br>7<br><i>x</i>                  | 6<br>8<br><i>x</i>                  | 6<br>9<br><i>x</i>                  | 6<br><i>a</i><br><i>x</i>                  | 6<br><i>b</i><br><i>x</i>                  | 6<br><i>c</i><br><i>x</i>                  | 6<br><i>d</i><br><i>x</i>                  | 6<br><i>e</i><br><i>x</i>                  | 6<br><i>f</i><br><i>x</i>                  |
| Row 2 | 6<br>5<br>0<br><i>x</i>             | 6<br>5<br>1<br><i>x</i> | 6<br>5<br>2<br><i>x</i>             | 6<br>5<br>3<br><i>x</i>             | 6<br>5<br>4<br><i>x</i>             | 6<br>5<br>5<br><i>x</i>             | 6<br>5<br>6<br><i>x</i>             | 6<br>5<br>7<br><i>x</i>             | 6<br>5<br>8<br><i>x</i>             | 6<br>5<br>9<br><i>x</i>             |  | 6<br>5<br><i>b</i><br><i>x</i>             | 6<br>5<br><i>c</i><br><i>x</i>             | 6<br>5<br><i>d</i><br><i>x</i>             | 6<br>5<br><i>e</i><br><i>x</i>             | 6<br>5<br><i>f</i><br><i>x</i>             |
| Row 3 | 6<br>5<br><i>a</i><br>0<br><i>x</i> |                         | 6<br>5<br><i>a</i><br>2<br><i>x</i> | 6<br>5<br><i>a</i><br>3<br><i>x</i> | 6<br>5<br><i>a</i><br>4<br><i>x</i> | 6<br>5<br><i>a</i><br>5<br><i>x</i> | 6<br>5<br><i>a</i><br>6<br><i>x</i> | 6<br>5<br><i>a</i><br>7<br><i>x</i> | 6<br>5<br><i>a</i><br>8<br><i>x</i> | 6<br>5<br><i>a</i><br>9<br><i>x</i> | 6<br>5<br><i>a</i><br><i>a</i><br><i>x</i> | 6<br>5<br><i>a</i><br><i>b</i><br><i>x</i> | 6<br>5<br><i>a</i><br><i>c</i><br><i>x</i> | 6<br>5<br><i>a</i><br><i>d</i><br><i>x</i> | 6<br>5<br><i>a</i><br><i>e</i><br><i>x</i> | 6<br>5<br><i>a</i><br><i>f</i><br><i>x</i> |

- Routing table of node with ID  $i = 65a1fcx's$ 
  - For each prefix  $p$  of  $i$ , and each digit  $d \in [0 \dots f]$ , has contact with ID prefix  $pd$


# Looking up objects in Pastry



- “Corrects” one digit at a time
  - Queries nodes with IDs: `d13da3`, `d4213f`, `d462ba`
  - Then use “leaf set” to find node with nearest ID to target

# Joining the system

- **Must know of one existing node in system**
  - Query it and other nodes to find node closest to your ID



- **Initialize leaf table from node closes to ID**
  - Will know almost complete leaf set for new node

# Initializing routing table

- **Can't initialize routing table from closest node**
  - E.g.,  $1\text{ffff}x$ 's closest node might be  $200000y$
- **But can fill up routing table from intermediate nodes**
  - Can use entire first row of first node contacted
  - Use second row of second node contacted, since same first digit as joining node
- **Once join procedure complete, can issue queries**
  - New node knows enough to route to and ID
- **But what about queries *to* IDs near new node?**

## Fixing up state for a join

- **New node must fix *other* nodes' routing tables as well as initialize its own**
- **For correctness, need to fix up neighbor's leaf sets**
  - Easy, node can contact them after initializing its own leaf set
  - If leaf sets correct, routing works, but could take  $O(N)$  hops
- **Updating other nodes' routing tables:**
  - Old routing tables either correct, or missing entry new node could fill
  - Automatically fill holes as side affect of lookups
  - New node sends its state to each node in its routing table
  - Nodes periodically query to try to fill holes in their tables

# Node failure

- **Nodes can fail without warning**
  - Other node's routing tables & leafs sets point to dead node
- **Routing table: Detect timeout, treat as empty slot**
  - Route to numerically closest available
  - Repair: Ask any node on same row for its contact  
Or ask any node below, since all will have correct prefix
- **Leaf sets: Node closest to target could be dead**
  - Need to find next closest
  - That's why leaf sets not just one neighbor ( $O(\log N)$ )
  - Easy to update leaf sets by contacting other nearby nodes

# How reliable is Pastry?

- **For correctness, only need leaf sets**
- **Assume independent node failures**
  - Each node fails with probability  $p$  in maintenance interval
  - Say leaf set contains  $L$  values
  - Probability of being cut off is  $p^L$
  - So for large  $N$ , if  $L \sim \log N$ , pretty good
- **Is independent failure a reasonable assumption?**
  - Good that  $\text{nodeID} = \text{MD5}(\text{IP Address})$
  - Proximity in ID space not correlated with physical proximity
  - But big network outages, synchronized renumbering correlated

# Locality

- **Lookup takes  $O(\log n)$  hops**
  - But hops could be long (NYC→Germany→LA)
- **Note: Many options for top levels of routing table**
  - Can chose *any* node with correct prefix
  - So pick nodes that are close to you to speed lookup
  - But makes it harder to assume independent failures
- **Continuously adjust routing table for locality**
  - Asks current entry for that entry's complete tables
  - Ping suitable nodes from other node's tables
  - Use them instead of current entry if ping says closer
- **No choice for leaf sets, though**



# Short routes property

- **Locality optimization helps recursive lookups**
  - New node will know of nodes close to it
  - Very good if triangle property holds (X close to Y and Y close to Z  $\implies$  X close to Z)
  - Often does hold, but not always
- **This is known as *short routes property***
  - Individual hops are lower latency
  - But less and less choice (lower node density) as you get close in ID space
  - So last few hops likely to be very long.
  - You don't *end up* close to the initiating node, just get there more quickly

# Scribe

- **Pastry can be used to form multicast trees**
  - Hash name of multicast group to get ID
  - Node closest to ID is *rendez-vous point* or root
  - To multicast a message, deliver it to RP, which sends it down the tree
- **Form multicast tree by routing JOIN msgs to ID**
  - Each node keeps track of groups + children for each group
  - On receipt of JOIN message, add sender to children
  - If child joins a new group, send join to parent (parent is just next hop towards ID)
- **Send just proceeds from RP to leaves**
  - Senders cache IP address of RP to save upwards routing
- **Leave protocol similar to join**

# Scribe locality

- **Short routes property helps multicast trees**
  - Towards leaves, parents are in high-levels of routing table
  - These are precisely the contacts with best locality
  - So often delivering messages to nearby nodes
  - Which may well reduce link stress (e.g., node 1abc $x$  at NYU will chose node 2def $y$  at NYU over farther nodes)
- **“Bottleneck remover” algorithm for overload**
  - Node may decide it is forwarding too many copies
  - Measures children & boots furthest away
  - Booted node effectively gets pushed down the tree

# Reliability & failure

- **Scribe sends messages over TCP**
  - But doesn't guarantee reliability
  - Nodes can crash and leave system abruptly
  - In fact, Scribe itself doesn't guarantee reliable delivery
- **Detect failures using heartbeat messages**
  - Each non-leaf node periodically sends heartbeat to children
  - Any multicast message serves as a heartbeat
  - So only need extra traffic when group quiescent
  - Upon timeout, route around failed node in Pastry
- **Must replicate root state in case root fails**
  - Typically replicated on 5 nearest nodes to ID

## **Reliable/ordered multicast**

- **Can build reliable/ordered multicast on Scribe**
- **Source assigns sequence number to each message**
  - Nodes do not send messages out of order
  - Wait for all previous messages before sending next
  - After fault+repair, you know what you are missing
- **Nodes buffer old messages**
  - Keep around for longer than detect+repair time
  - So when you repair, can request messages you missed

# Splitstream

- **Problem: Scribe makes uneven use of resources**
- **In fully-balanced tree w. height  $h$ , fanout  $f$** 
  - $f^h$  leaf nodes consume no upstream b/w
  - $(f^h - 1)/(f - 1)$  internal nodes consume  $f \times$  stream b/w
  - E.g., with  $f = 16$ ,  $< 10\%$  of nodes carry forwarding load!
- **Better approach: Stripe data over a *forest* of trees**
  - Each node is leaf in some, internal in others
  - Could round-robin packets down multicast trees
  - Or could stripe at the bit level
  - One tree could be parity bit, to survive a failure

## Interior-node-disjoint trees


- **Want to avoid a failure affecting multiple streams**
  - E.g., say node  $n$  is your ancestor in multiple trees
  - If  $n$  fails, you lose multiple streams (so parity won't help)
- **Solution: Each node is interior in only one tree**
  - Say digits are in base 16
  - Can achieve by having 16 trees, each with a group ID that starts with a different digit
  - Can only be interior node if group ID and you have at least one-digit prefix in common

# CFS

- **Another application of P2P systems**
- **Idea: Replicate widely stored data in DHT**
  - E.g., Linux distribution
  - Care a lot about data integrity—no tampering!
- **CFS – cooperative file system is P2P file system**
  - Read-mostly file system
  - Publish operation breaks into blocks
  - Spreads chunks all around DHT
  - Digitally sign entire file system for integrity

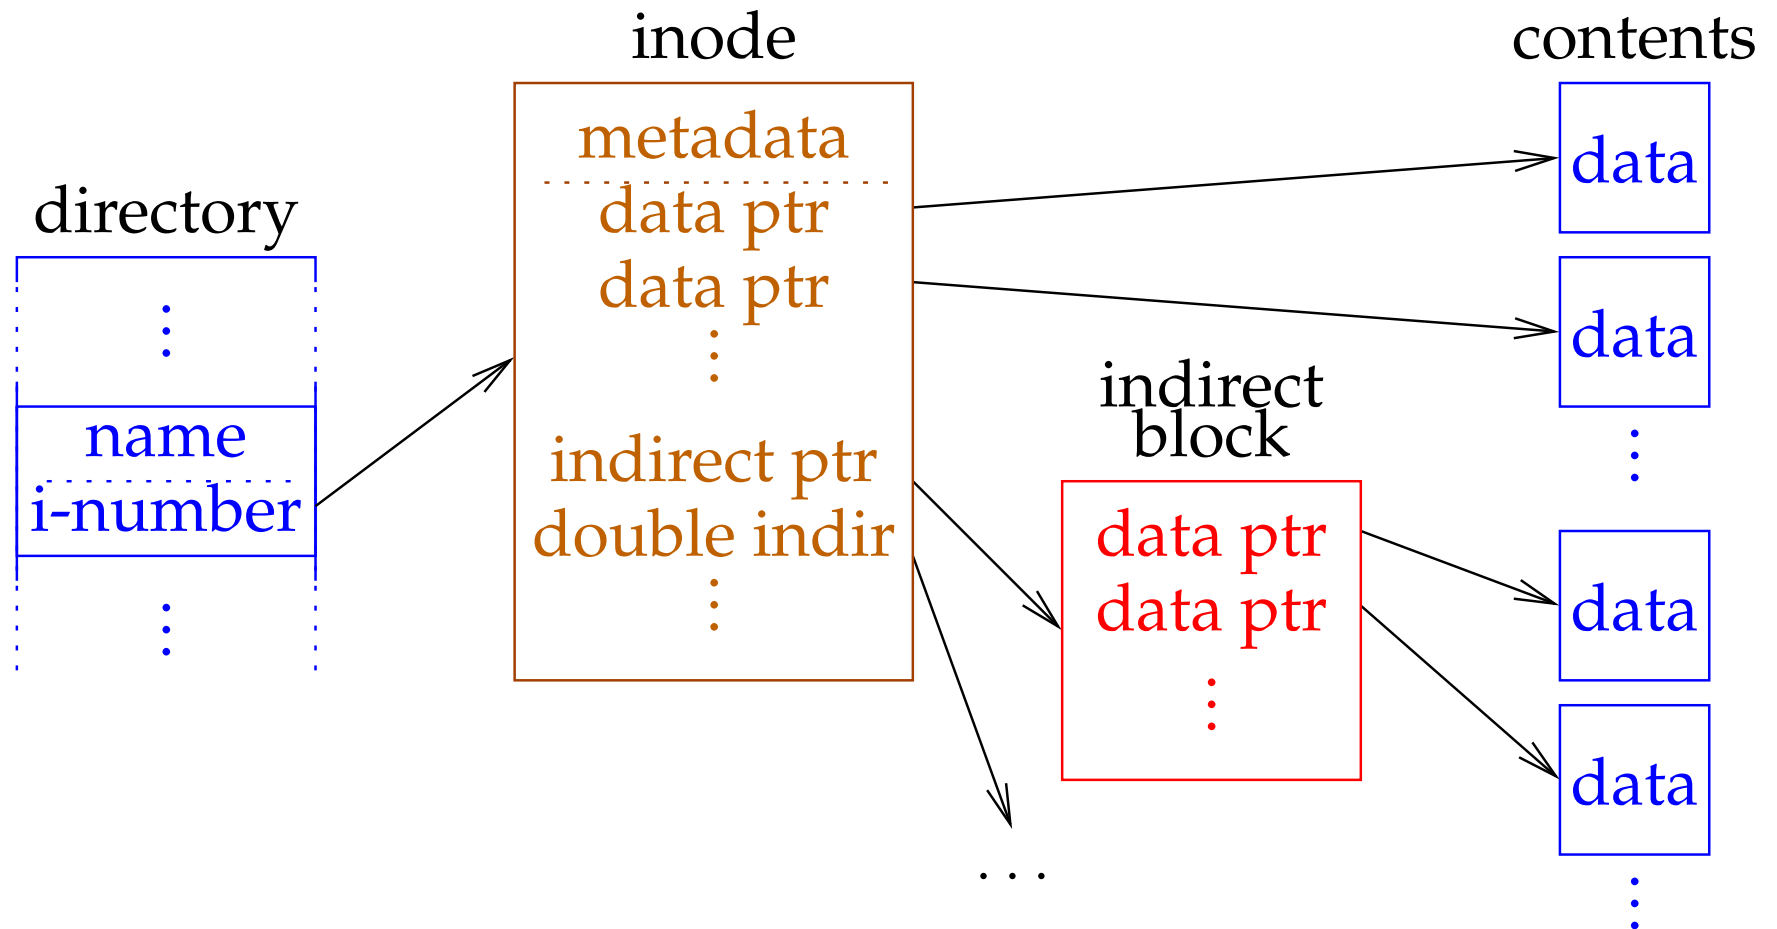


## Example: Publishing 2 blocks of data



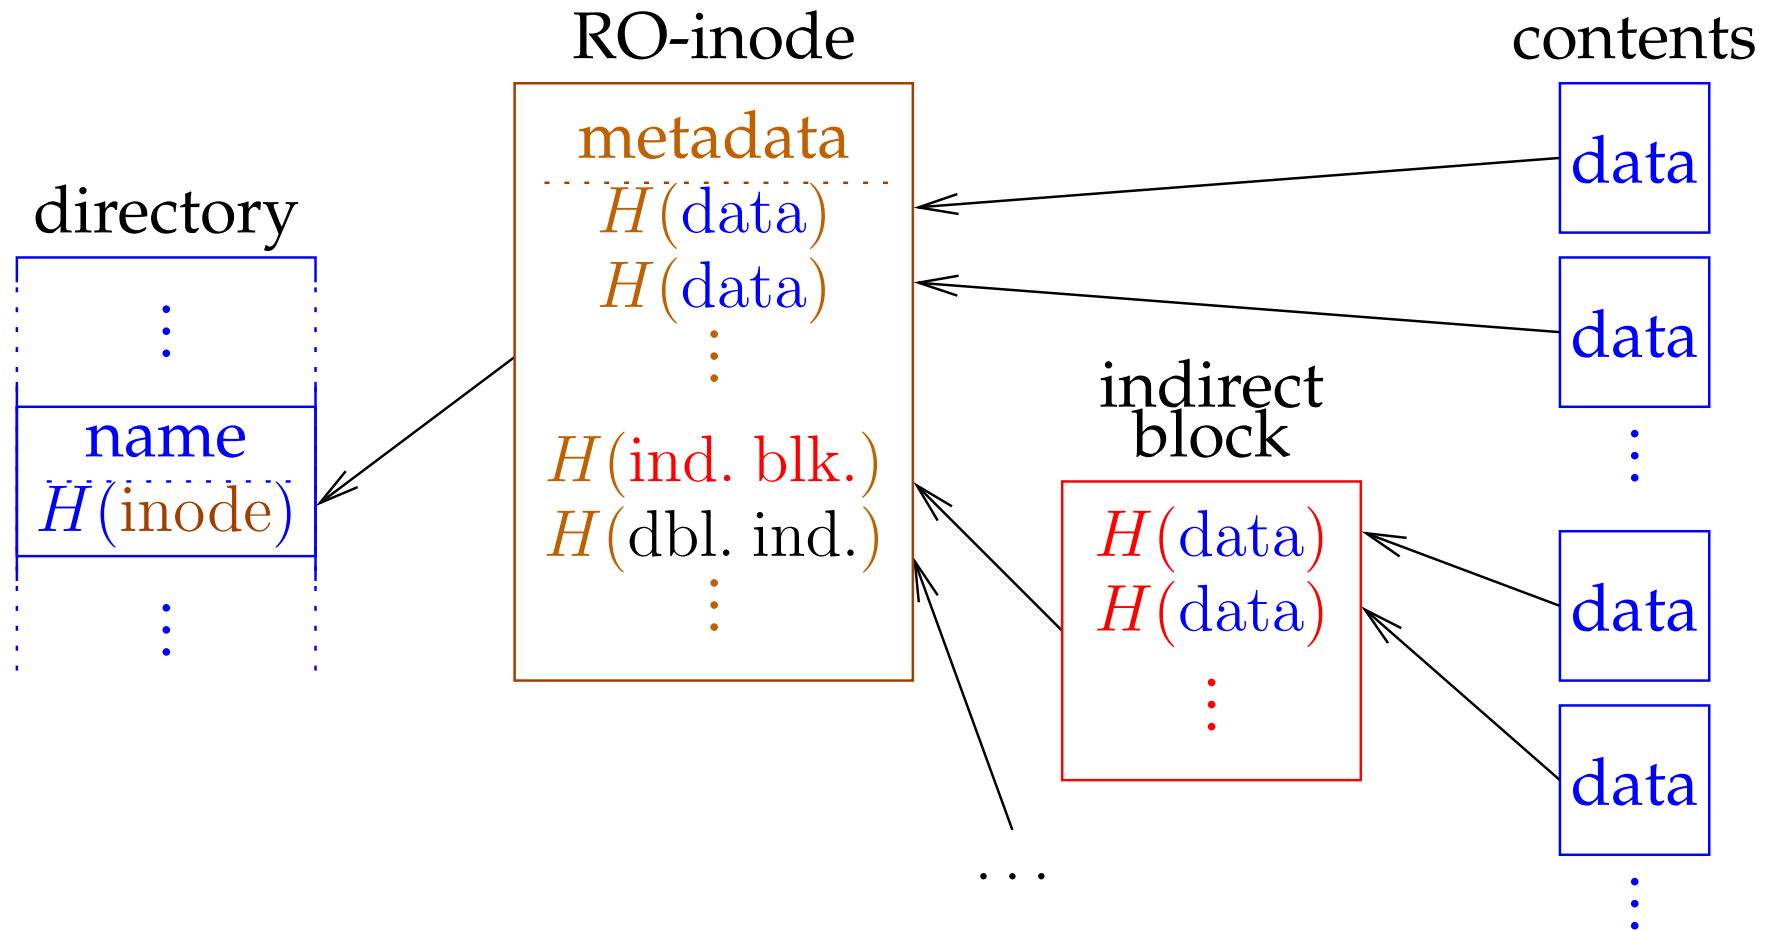
- **Digitally sign version & hashes of blocks**
  - Can verify one block without having the other
  - Two blocks must come from same version of file
- **Generalize technique to an entire file system**

# Traditional FS data structures



- In database arbitrary key can replace disk location

# CFS data structures



- Index all data & metadata with cryptographic hash

# CFS scalability & reliability

- **CFS built on Chord not Pastry, but ideas similar**
- **Blocks must be replicated for reliability**
  - Easy: Store each item at  $k$  successor nodes around circle
- **Blocks must be replicated for scalability**
  - E.g., Imagine everybody reads the same block
  - Don't want to overload poor successor node
- **Solution: Store blocks along the lookup path**
  - Suppose you are looking up block  $B$  on node  $n_0$
  - You may traverse nodes  $n_3, n_2, n_1, n_0$  to get  $B$
  - Now store  $B$  on  $n_1$
  - Next lookup that converges at  $n_1$  will store on prev, etc.